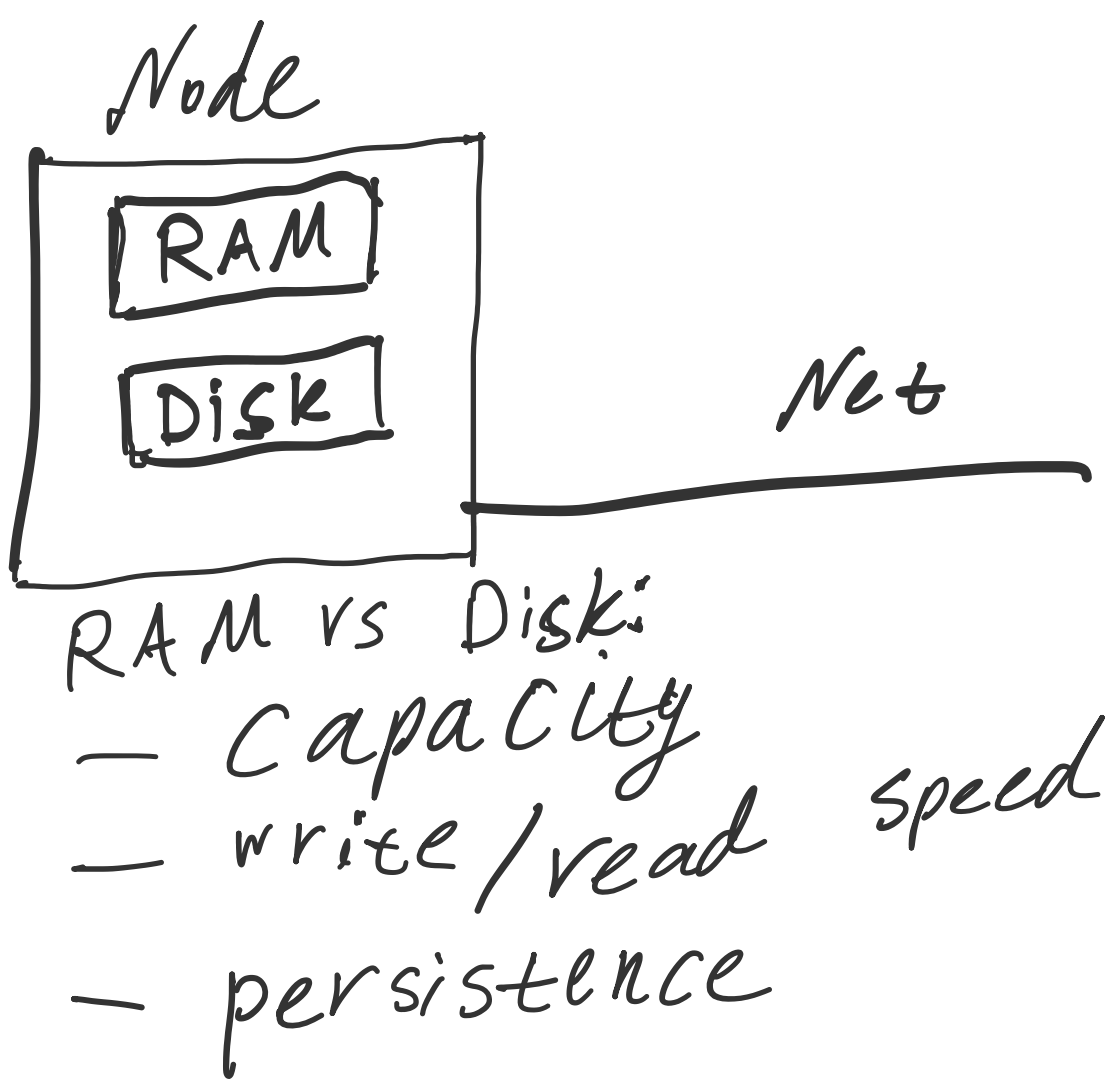


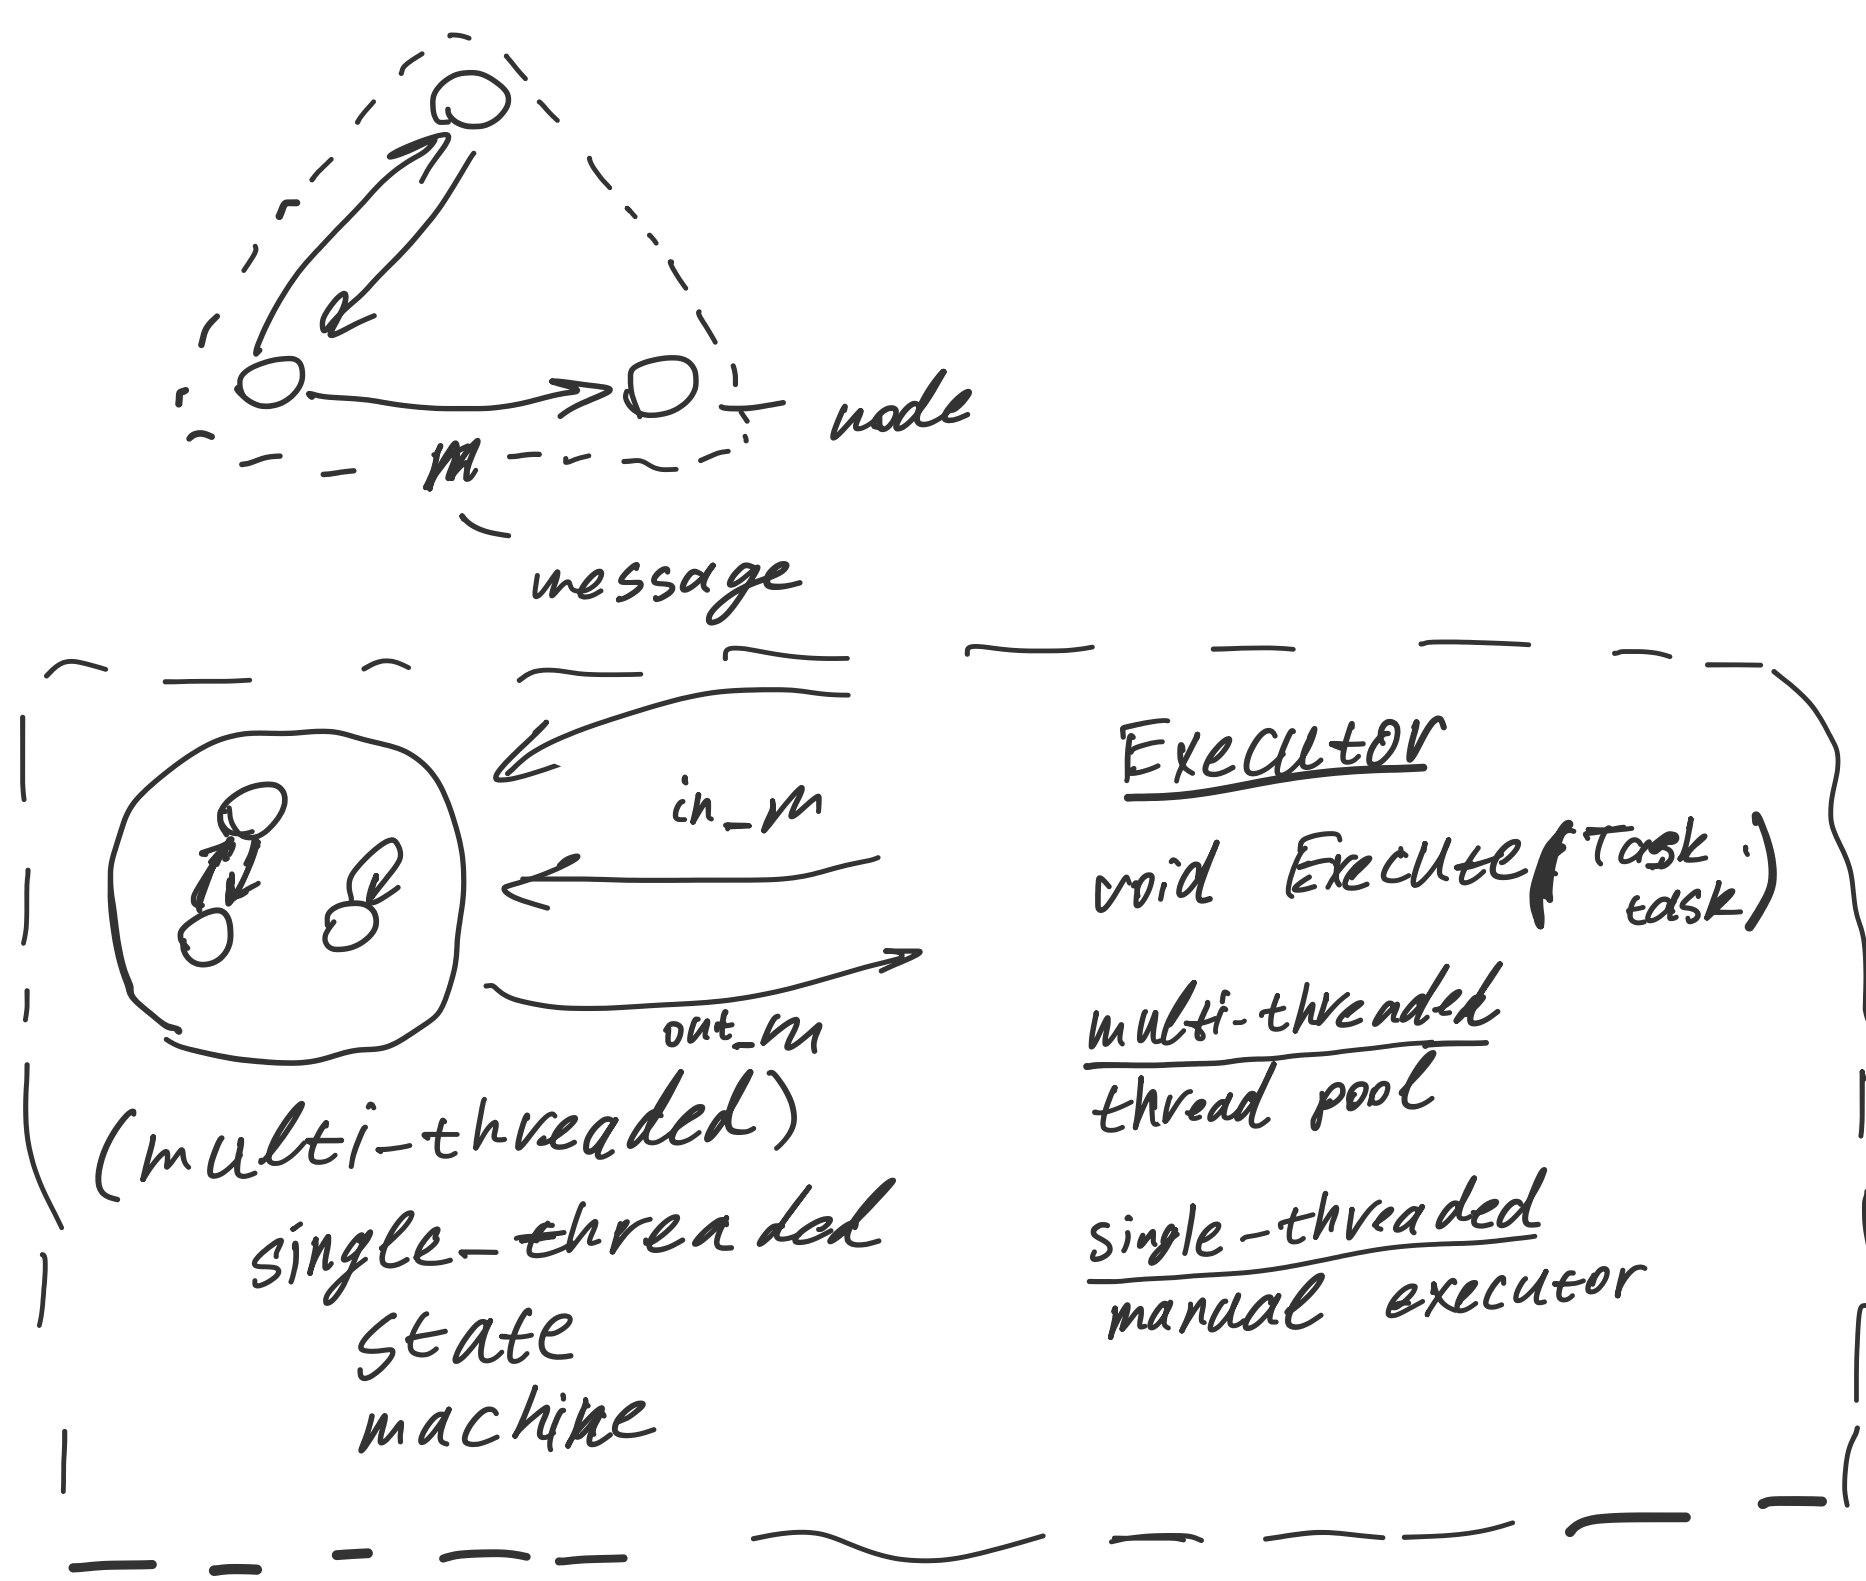


- filesystem
- database
- key value storage (k,v)
- queue / message broker
- coordination service

Why distributed?

- Scalability
- Fault-tolerance
- Byzantine faults\*




Node:

- 1) Crash
- 2) Restart
- 3) Byzantine faults\*

Net:


- 1) Link break
- 2) Network partition
- 3) Byzantine faults\*

Network models



• sync

• async




Partitioning (Sharding)


Replication

shard 0	shard 1	
[abc]	[123]	replica 0
[abc]	[123]	replica 1


shard 0



shard 1




+ shard 2



shard manager

for load balancing



$k \rightarrow v$

$k_1 \rightarrow k_2 \rightarrow v$   
read/write skew

$\text{hash}(k^*) \% 2 = 1$

$\text{hash}(k^*) \% 3 = 2$

